

# Hibernate Summary

Author : Ankit Panda  
Mail : akpanda.ds@gmail.com  
Sample implementation

<https://github.com/akpandads/HibernateImpl>  
<https://github.com/akpandads/AdvanceJpa>

## Accessing DB

ORM model which maps tables to classes and fields to variables  
Use inmemory database like H2 for demo purposes  
JPA is the interface while hibernate is the implementation

If using spring boot, autoconfiguration feature configures all the database connections  
There are different ways of retrieving data from DB

### 1. Spring JDBC template

- Autowired by spring boot
- takes parameter
- Advantageous over regular JDBC , no boiler plate code required to map rows and manage transactions
- Uses a BeanPropertyRowMapper to map the result to a entity class

```
jdbcTemplate.query("select * from person", new BeanPropertyRowMapper<Person>(Person.class));  
---- direct mapping  
jdbcTemplate.queryForObject("select * from person where id=?", new Object[]{id}, new  
BeanPropertyRowMapper<Person>(Person.class)); --parameter passing via Object Array
```

class PersonRowMapper implements RowMapper < Person > { -----Inner class in the  
entity object to for custom row mapper. Note the default can also be used if the table columns and  
class variables name matches

```
@Override  
public Person mapRow(ResultSet resultSet, int i) throws SQLException {  
    Person person = new Person();  
    person.setId(resultSet.getInt("id"));  
    person.setName(resultSet.getString("name"));  
    person.setLocation(resultSet.getString("location"));  
    person.setBirthDate(resultSet.getDate("birth_date"));  
    return person;  
}  
}
```

### 2. Entity Manager

- Persistence context provides a bean entity manager
- Persistence context maintains all the beans that are associated with the rows

```
@PersistenceContext  
EntityManager entityManager;
```

```
entityManager.find(Person.class, id);  
entityManager.merge(person);
```

TypedQuery <Person> typedQuery = entityManager.createNamedQuery("find\_all\_persons", Person.class); ----- the  
typed query name is defined at the entity class. here just the name of the query is used  
return typedQuery.getResultList(); ■

### 3. Spring Data Jpa

- No need for creating Entity Manager and repository
- Implement the Jpa Repository or Crud Repository and use available functions
- Not suitable for complex queries or the one which require performance optimization
- Supports out of box sorting and pagination functionalities
- inbuilt methods like find(), findAll(), findBy<ElementName>() are available

# JPA annotations

We will look at some common JPA annotations and methods on this section

**@Entity** to Denote Entity class

**@Table(name= <table name>)** used if the entity name is different from table name

**@Column(name =<name>)** denote mapping between column and entity variable

**@NamedQuery(name="get\_all\_courses",query="select c from Course c")** = Query to be used with name. Defined at top of entity class

**@UpdateTimeStamp** = java type can be mapped to LocalDateTime. Used to store timestamp of last updation on the row

**@CreationTimeStamp** = java type can be mapped to LocalDateTime. Used to store timestamp of creation on the row

**@Id** = Denote primary key

**@GeneratedValue** = denote that the value is generated from a generator

**@Transactional** = When used on a method or class, all the db operations on that method is condered as one transaction

**@DirtiesContext** = Used on test methods usually to denote that the method is making changes on the persitence context. And the changes made as part of this method will be reverted back

**@Enumerated** = Use this whenever there are enums being used. This will ensure that string values are stored in DB and not the positional values of the enums

When are the methods called

If the fetch type is eager, then all the associated data are called at once.

if fetch type is s lazy the queries are fired when the associated data are fetched

Hibernate usually holds on persisting data till the last moment. All the changes done on entity objects are stored in the persistence context.

At end of the method all perstent context objects are persisted in the db, that is yupdates are executed

If to force the changes to DB at regular intervals, the flush() can be used to force execute the the updates , instead of waiting for hibernate

The detach() is used to detach an object from the persistence context. once detached, the object any changes to the object are not stored in DB and it is not associated with any transactions

The clear() clears all the changes done to the entity object and not yet persisted in DB. In other words, the state of the enity is synced to the values in DB. All local changes in persistence context are lost

## Relationships

there are 4 different types of rellationships that can be represented.

OneToOne, OneToMany, ManyToOne, ManyToMany

\*ToOne are Eager fetch by default

\*toMany are Lazy fetch by default

Bitdirectional Mapping, is where there exists a relation on boths sides from two entity classes . Changes in one entity object will make changes in the associated entity oobject of the other class

Unidirectiona Mapping is when there is only one way relationship between 2 entities. Chenges in the owning entity reflect in the owned entoty but not the other way round.l

**@OneToOne**

Consider 2 tables Student and passport. One student can have only one passport

Entity Student definition

```
@OneToOne(fetch = FetchType.LAZY) --no mappedBy denotes owning side. here stident is the owning side
private Passport passport;
```

Entity Passport Definition

```
@OneToOne(fetch = FetchType.LAZY, mappedBy = "passport") --mappedBy is not present on owning side of the
relationship
private Student student;
```

Save Student with passport

```

public void saveStudentWothPassport(){
    Passport passport = new Passport("G176251276");
    entityManager.persist(passport);
    Student student = new Student("Mike");
    student.setPassport(passport); entityManager.persist(student); ---note both the entities need to be oersisted
separately
}

```

## @ManyToOne

Many reviews can be associated to one course

Review Entity Dfinition

```

@ManyToOne
private Course course;

```

## @OneToMany

Each course entity can have many reviews

Entity Course definition

```

@OneToMany(mappedBy = "course")
private List <Review> reviews = new ArrayList<>();

```

## @ManyToMany

One course can have many students

```

@ManyToMany(mappedBy = "courseList")
private List <Student> students = new ArrayList<>();

```

One student can be in many courses

```

@ManyToMany
@JoinTable(name = "STUDENT_COURSE", joinColumns = @JoinColumn(name = "STUDENT_ID"),
           inverseJoinColumns = @JoinColumn(name = "COURSE_ID") )
private List <Course> courseList = new ArrayList<>();

```

The joinTable creates another table which contains just the student id and course id. Each entry will correspond to a student id who is enrolled on the course

Inserting a row with student and course

```

public void enterStudentAndCourse(Student student,Course course){
    entityManager.persist(student);
    entityManager.persist(course);
    course.getStudents().add(student); ---- this is required to make the mappinf from course to student
    student.getCourseList().add(course); ----- this is required to make the mapping from student to course
}


```

# Inheritance Strategy

In this section we will discuss the 4 inheritance stratgy and Embeddable concpet

Consider the scenario where there is an abstract class Employee

2 classes FullTimeEmployee and PartTimeEmployee extend the Employee class. We will look at how this can be represented in hibernate and DB



### 1. InheritanceType.SINGLE\_TABLE

The variables from Employee, PartTimeEmployee and FullTimeEmployee are combined into one table and all the variables are present as columns in a single table

This type is great for performance because only one table is present but poses a high risk for data integrity. A single method on the superclass can be used to retrieve data. Since only one table, no joins are involved. The columns which do not correspond to a particular table, are filled up as null for that particular row. We can use a `@DiscriminatorColumn` to create a column to determine the type of a particular row.

Id	name	salary	hourly
1	x	null	12
2	y	1200	null

### 2. InheritanceType.TABLE\_PER\_CLASS

A separate table per class is created containing only the common columns of Employee and the specific columns of the subclass.

id	name	salary

FullTimeEmployee Table

id	name	salary

Part Time Employee

### 3. InheritanceType.JOINED

This is a great model where data integrity is more important than performance.

3 tables will be created. One for each type of the class. And the id of the superclass will act as a foreign key in the subclass tables.

id	name

Employee table  
Employee Table

salary	id (this the same entry as that of employee table)

FullTimeEmployee Table

salary	id (this the same entry as that of employee table)

Part Time

### 4. MappedSuperClass

Mapped super class removes all the relation between the super class and subclass while forming table.

A separate method needs to be written for accessing data from each of the subtables as the super class object will not be considered as the generic one.

The tables created will be similar to table per class strategy with the difference that the abstract or super class is not an entity.

So polymorphic queries cannot be written.

### Embeddable and Embedded

A composite object is one which contains another object within itself.

The Composite class needs to have the `@Embedded` on the enclosing member variable.

The smaller class needs to have the `@Embeddable` annotation on it to denote that this can be embedded on a composite class.

In the DB representation, all the attributes of Embeddable class are expanded as represented as columns, within the composite class table.

## Transaction Management and Caching

Transaction Management is made easier with the `@Transactional` annotation provided by Spring.

But there are a lot of issues that need to be handled while managing a transaction. Each transaction must adhere to the ACID properties to maintain data consistency

**A- Atomicity** : the whole transaction be considered as a single unit. Either it should all persist or rolled back  
**C-Consistency** : Successive reads within a transaction must return the same number of rows and same value  
**I-Isolation** : Any number of transactions operating parallelly must not affect the other  
**D-Durable** : A transaction once successful must be persisted permanently

### **Problems in transactions**

**Dirty Read** : 2 simultaneous transactions. Transaction 1 changes value of x. Transaction 2 reads x. Transaction 1 fails and reverts x. But transaction 2 updates on the modified value. Basically a transaction reads the value before it is persisted by another transaction

**Non Repeatable read** : Transaction 1 reads x. Transaction 2 updates x while transaction 1 is still in progress. Now transaction 1 reads the same value and gets a different value for x

**Phantom Read**: Transaction 1 does a select query and gets n number of records. Transaction 2 inserts/deletes some record while transaction 1 is in progress. Transaction 1 again does a select and gets another set of results 'm' different than initial 'n'

Isolation levels

read uncommitted : solves none of the issues

read committed : solves dirty read. Lock on all the values changed while a transaction in progress

repeatable read: solves dirty read and non-repeatable read. lock on entire row

Serializable : solves all the problems. But performance issue. Everything locked that matches with query condition

### **CACHING**

2 levels of cache. L1 and L2.

Level 1 is bounded by each transaction. Once a value is fetched within a transaction boundary this is not again queried but fetched from the Persistence Context

Level 2 cache spans across multiple transactions and this needs to be specified in code what all data needs to be in level 2 cache.

@Cacheable to be declared on the entity which needs to be cached

```
spring.jpa.properties.hibernate.cache.use_second_level_cache=true
```

```
spring.jpa.properties.hibernate.cache.region.factory_class=org.hibernate.cache.ehcache.internal.EhcacheRegionFactory
```

```
spring.jpa.properties.javax.persistence.sharedCache.mode=ENABLE_SELECTIVE
```

The above 3 show the properties. EhCache is one of the L2 cache libraries. ENABLE\_SELECTIVE signifies by default none of the entities will be cached but we can select that all needs to be cached

The opposite is DISABLE\_SELECTIVE